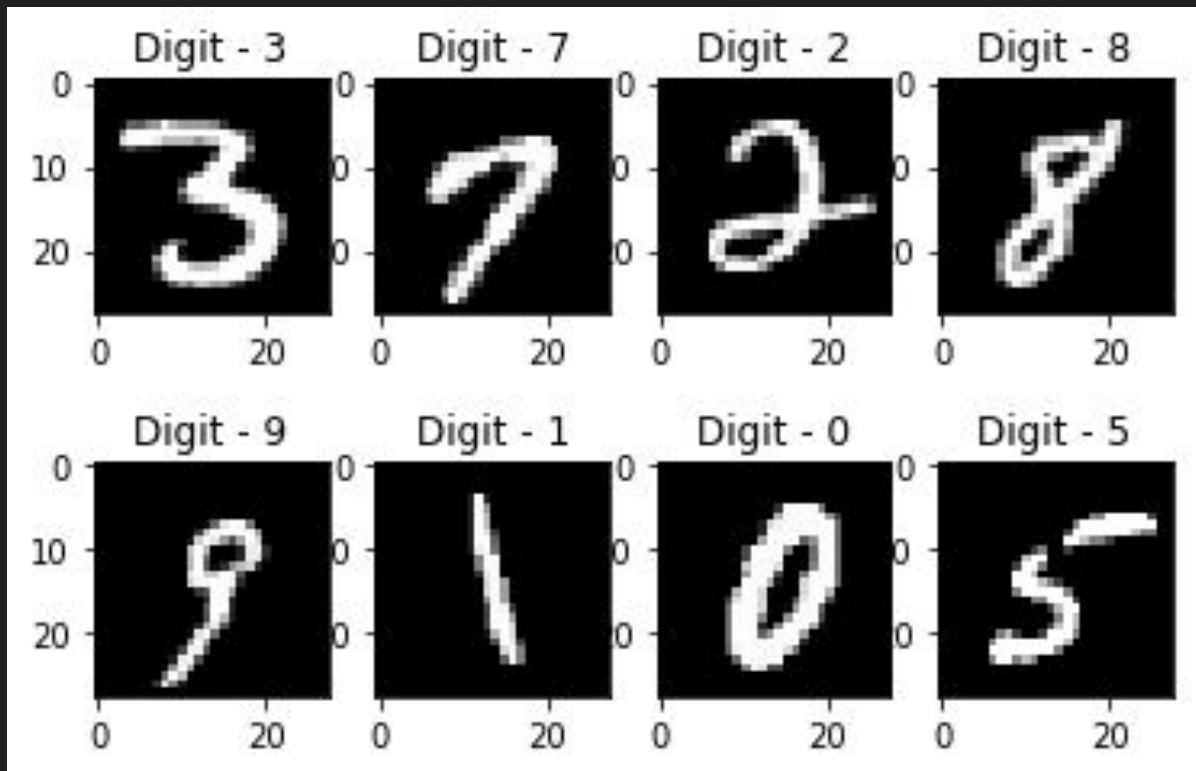


1 - Project Description


Introduction: FPGA

Accelerated Handwriting Digit Recognition



Adapted from [1]

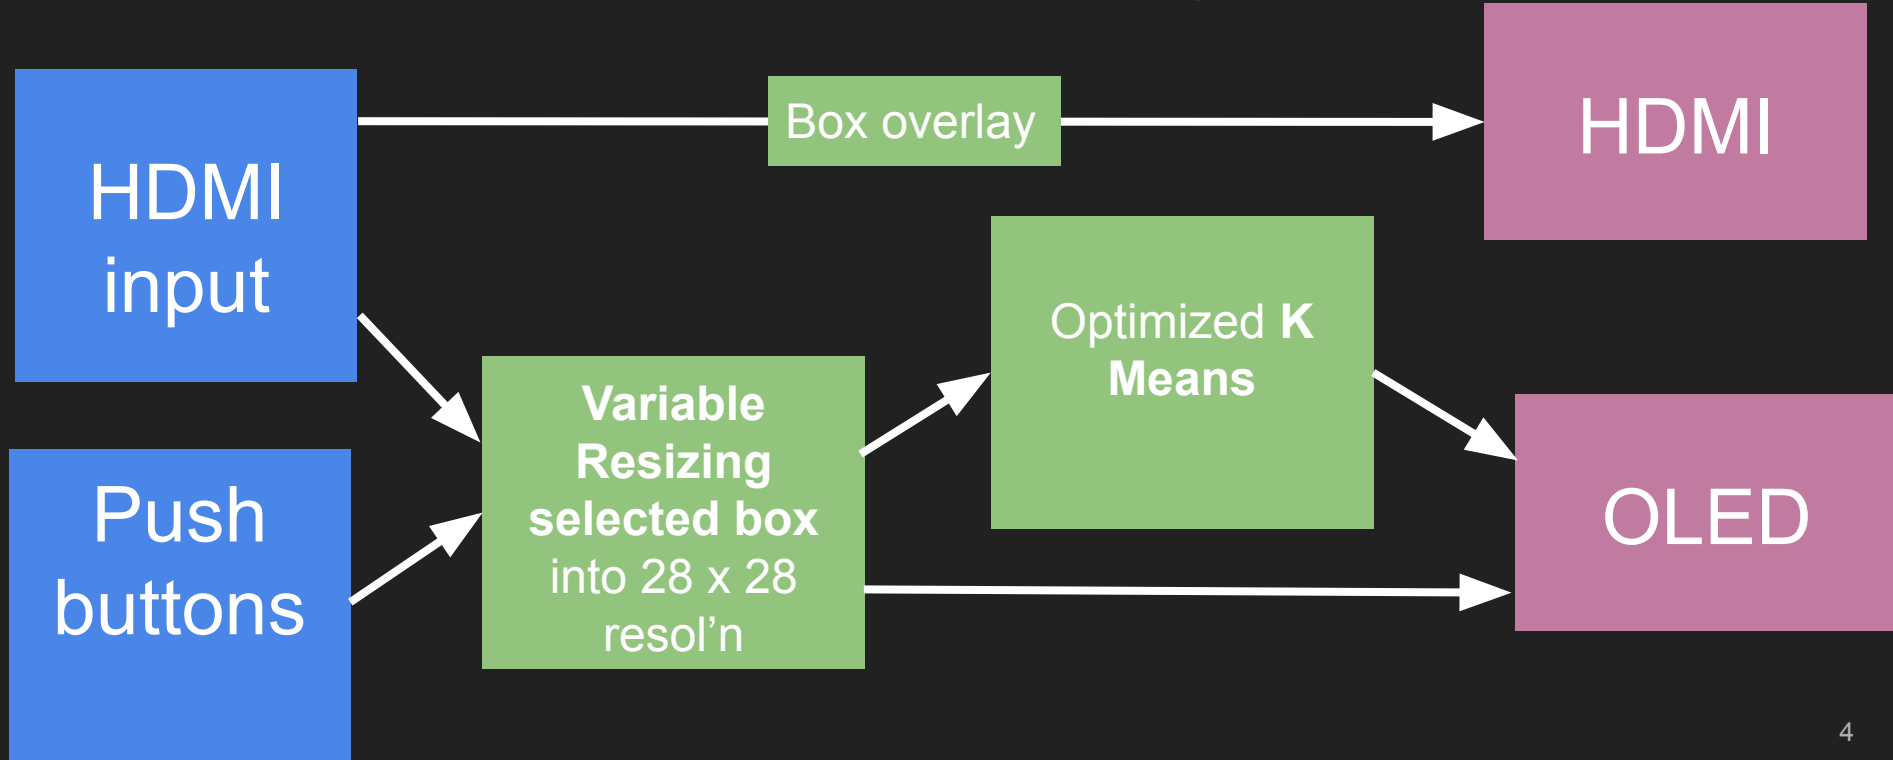
Neural Classifier: **KNN/K-Means**



Input

Processing

Output



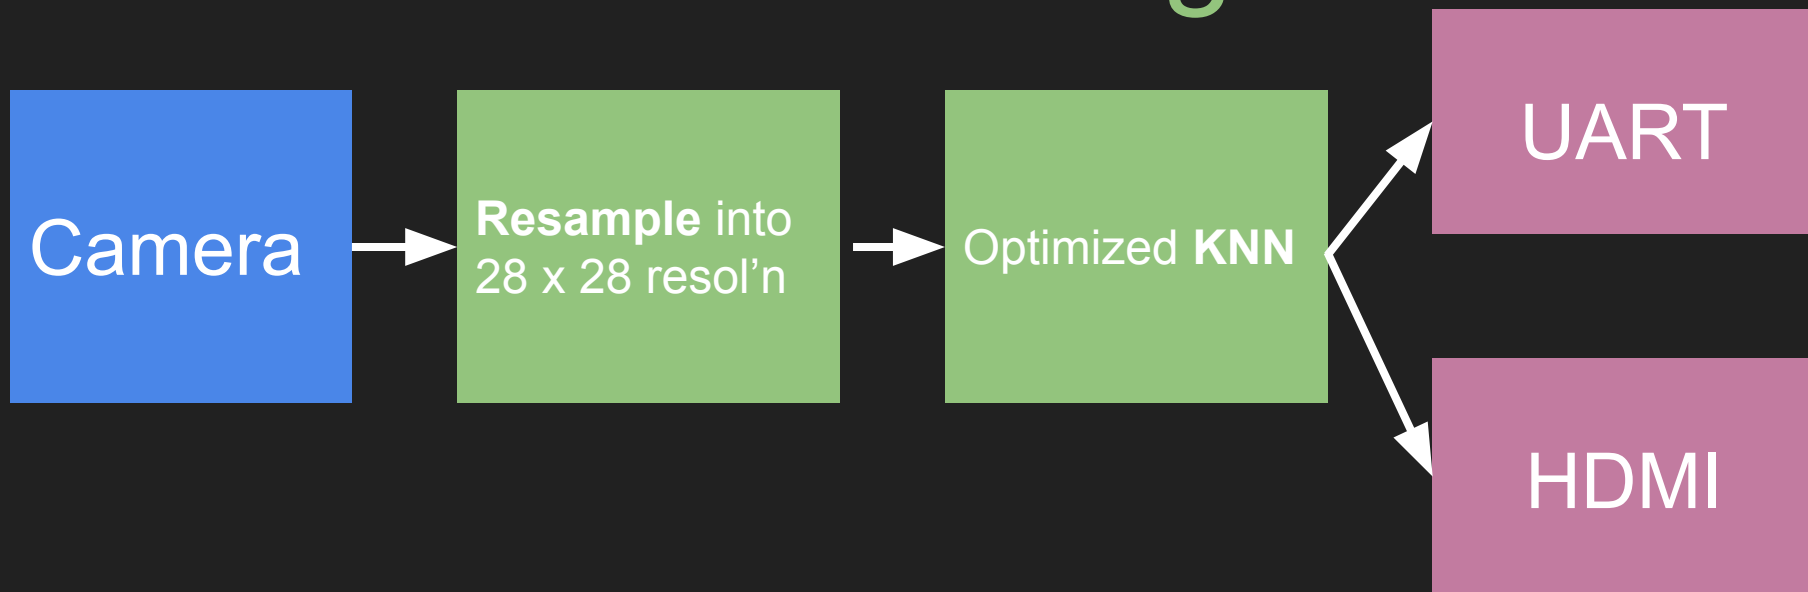
2 - Initial Goals

Initial Concept

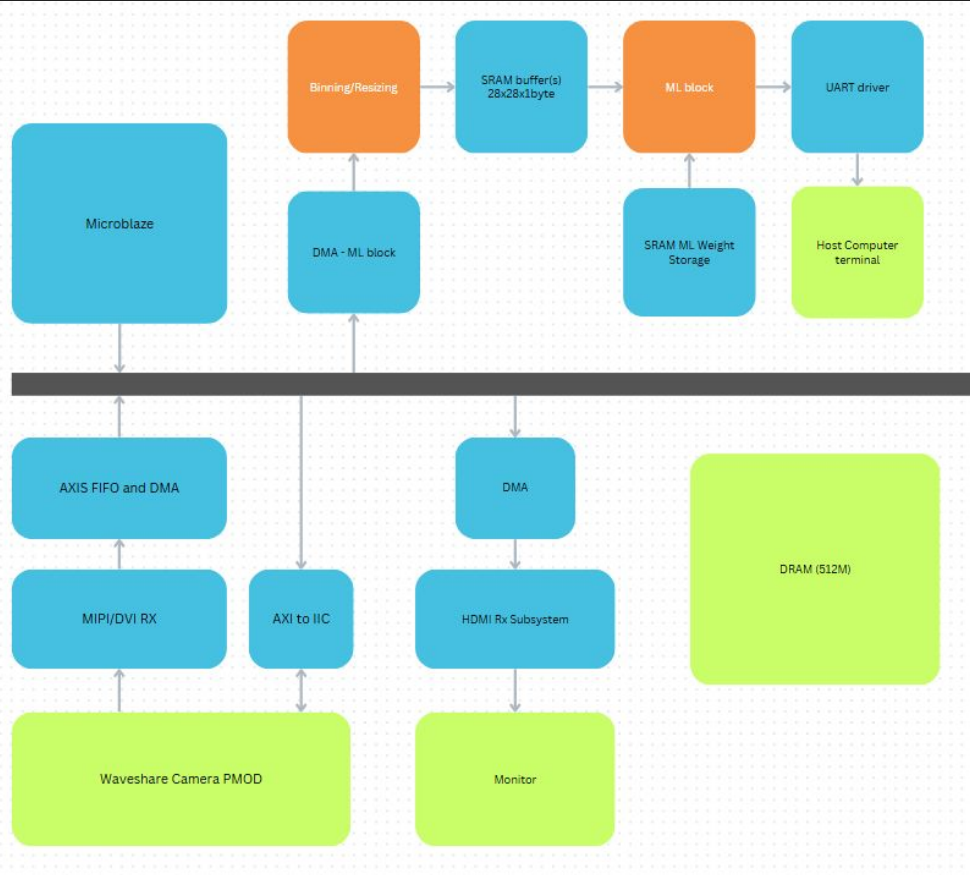
Input

Processing

Output



Initial Top Level Block Diagram



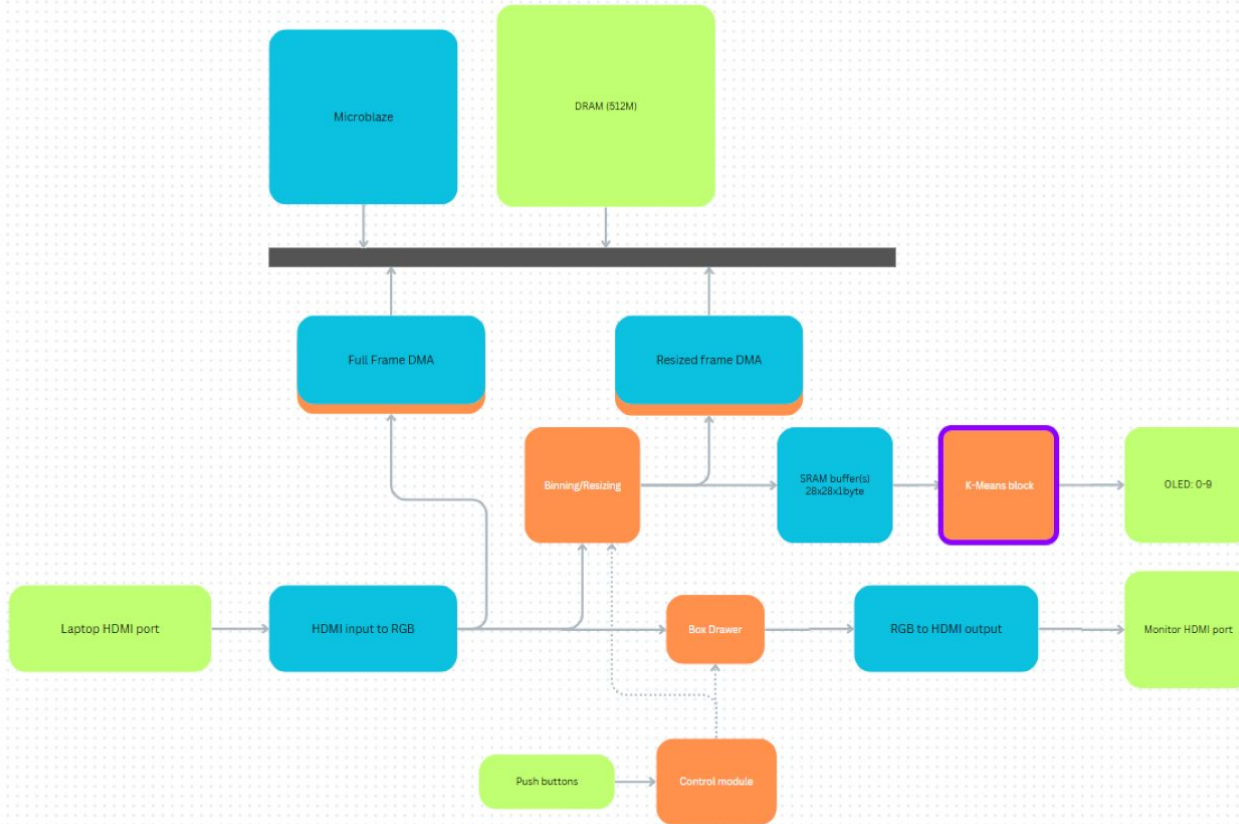
Main Changes

1. HDMI I/O not Camera
2. Variable input frame sizes
3. K-means using HLS
4. OLED to show results
5. Stream video instead of using

DDR + AXI

3 - Final Project - high level block diagram

Refined Top Level Block Diagram



Vivado IP

I/O devices

IP blocks

Vivado HLS

List of Custom IPs used

Packaged IPs

- K-Means classifier - fast + slow versions
- Resizing block
- Box controlling module
- Box drawer

Helper IP

- SW triggerable Video-2-MM-2-KMeans FSM
 - (controls interface between the K-Means + resizer)
- SW triggerable Video-2-AXIS (save 1x frame)
 - (controls interfaces with Sw)


SW IP

- K-Means in SW
- Resizing in SW
- Event loop + FPS counter


4 - Final design - Individual components

(1) Custom IP: K-Means

Input and output



IP creation



Utilization of K-means

Resource	Estimation	Available	Utilization %
LUT	23350	134600	17.35
LUTRAM	4	46200	0.01
FF	16459	269200	6.11
BRAM	784	365	214.79
DSP	740	740	100.00
IO	77	285	27.02
BUFG	1	32	3.13


Neural Classifier Performance Comparison

Classifier type	Accuracy on MNIST	Storage requirement	Computational Requirement (per inference)
Full KNN	97.05%	45 MiB	47M MACs
Reduced KNN	81.74%	784 KiB	802K MACs
K means	82.05%	7.66 KiB	7.86K MACs

Atrix 7 BRam capacity: 13 MBits

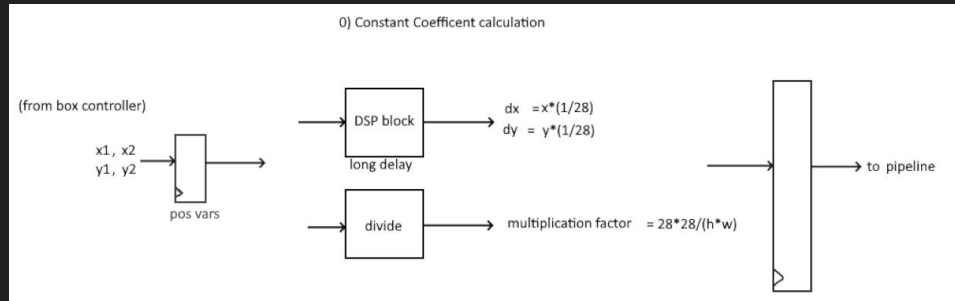
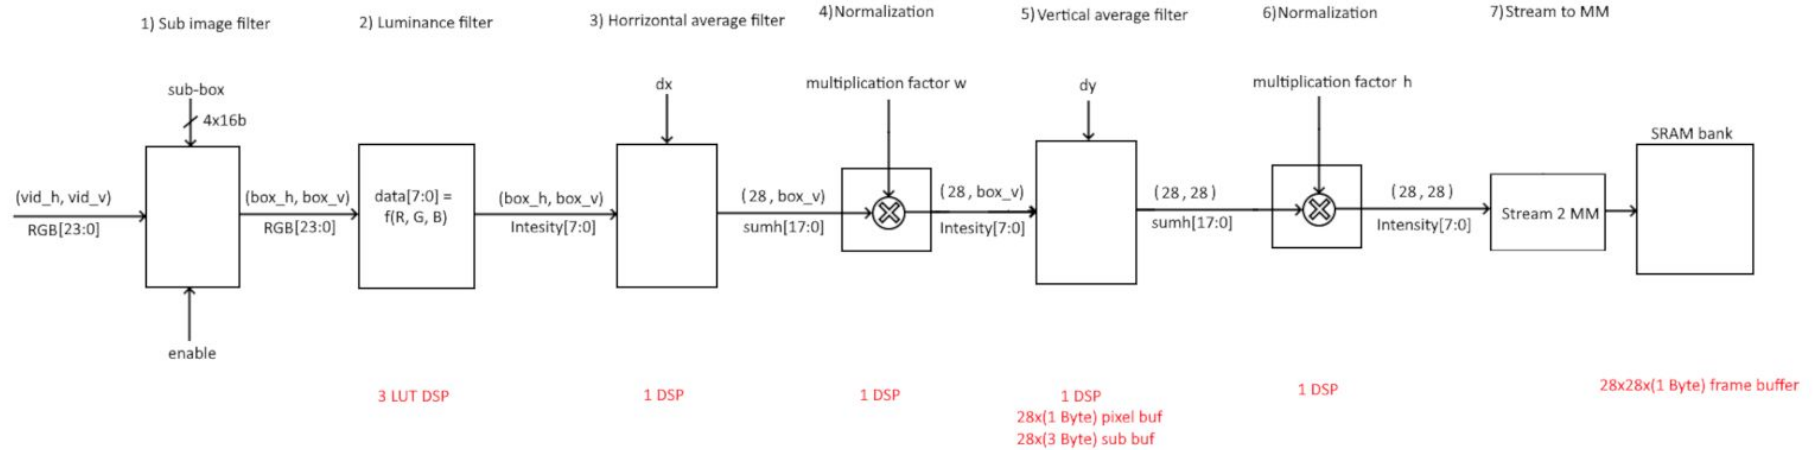
Atrix 7 DRAM capacity: 512MB

(2) Variable resizing

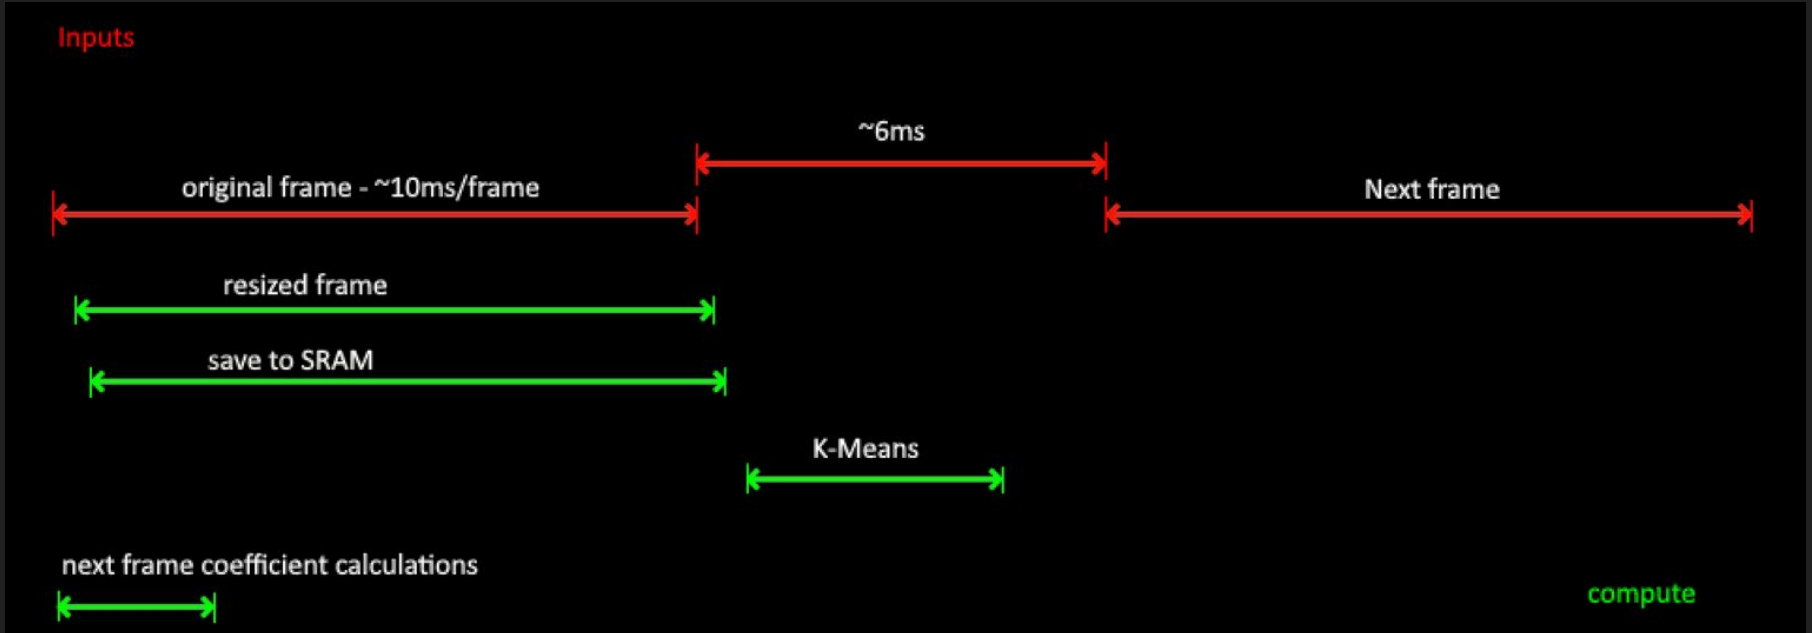


- Up sampling
- Down sampling
- Luminance
- Minimal HW resources and Latency

Resizing 7-Stage Pipeline



Timing diagram - Resizing IP



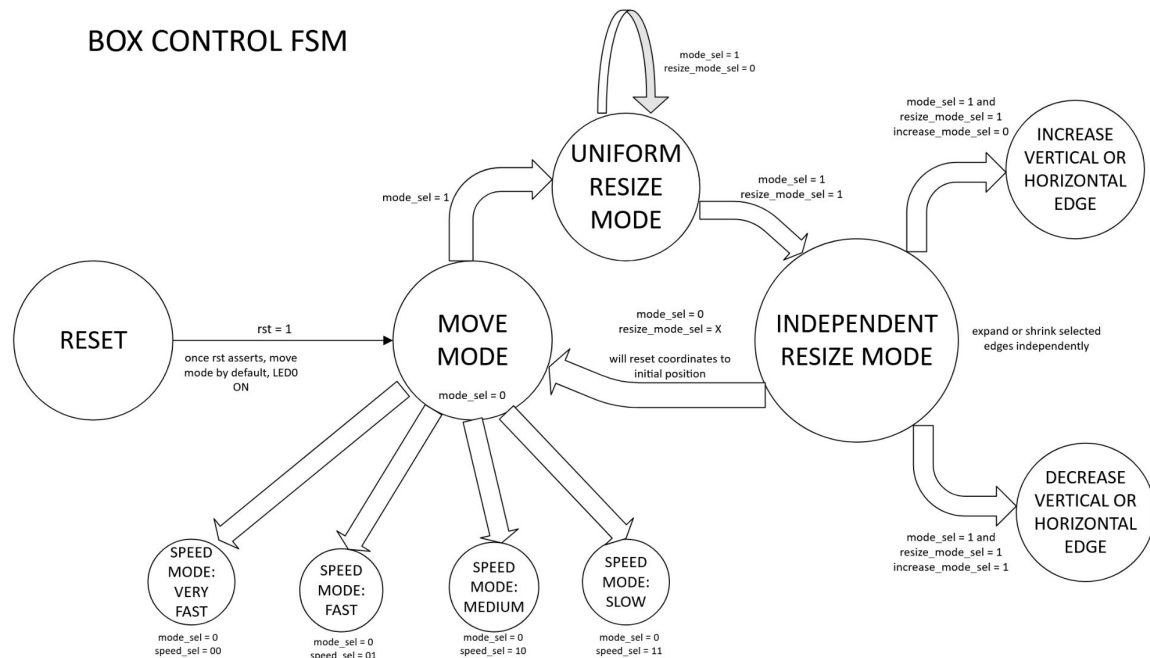
Utilization of Resizing IP

Resource	Utilization	Available	Utilization %
LUT	1153	133800	0.86
LUTRAM	48	46200	0.10
FF	1658	269200	0.62
DSP	4	740	0.54
IO	106	285	37.19
BUFG	1	32	3.13

(3) Box Control

- Controls the movement and resizing of the bounding box of the image frame
- 2 modes:
 - Movement: a) fast b) slow
 - Resizing: a) uniform b) independent edge resizing
- Moderate complexity
 - Edge-sensitive (rising edge detection)
 - Time-sensitive (delay-based movement control) functionalities
- Speed selection mode can impact responsiveness and performance

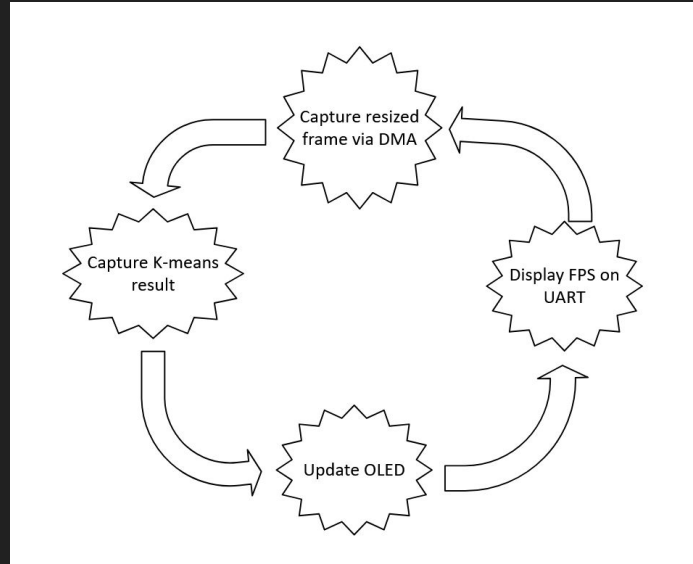
BOX CONTROL FSM



Utilization of Box Control

Resource	Utilization	Available	Utilization %
LUT	232	63400	0.37
FF	95	126800	0.07
IO	78	210	37.14
BUFG	1	32	3.13

(4) Software




(additional)

- `capture_full_frame()`
- `resize_image()`
- `infer_on_saved()`

`print_ascii_art()`
`print_saved_bitmap()`

(5) Data flow + supporting blocks

- Resize/Binning Module
- HDMI to RGB Block
- RGB to HDMI Block



5 - Difficulty Score

Project Complexity Points

Binning/Resizing Module IP (HW + SW)	1.0
Knn-Means Classifier Packaged IP (HW + SW)	1.0
Box controller IP	0.5
OLED implementation	0.75
HDMI output	1.0
HDMI input	1.0
Implement performance monitoring & Meaningful visualization statistics (fps)	0.25
Total	~5.5

6 - How would we improve with more time?

Possible Improvements

1. More accurate classification

- Use more centroids (i.e. >1) and parallelize distance calculations with centroids
- Preprocessing images (e.g. deskew, center, remove artifacts)
- Different distance metric
- Feature extraction
- (lenet)

2. Ping-pong buffering to increase throughput

3. Support for multi-character recognition and non-digit recognition

7 - The Demo

Demo Modes

1. SW resize, SW classification, and output to OLED
 - a. 1.62 fps (only at low sized sub images)
2. HW resize, SW classification, and output to OLED
 - a. 19.81 fps
3. HW resize, HW classification, and output to OLED
 - a. 29.72 fps
4. HW resize, HW classification, only value pushed to OLED (60fps)
 - a. Full 60fps

Results

Function	SW Latency (on microblaze)	HW Latency	Speedup
K-Means	15.7ms	235.4 us / 4.09 us (optimized)	99%
Update OLED with text / 28x28 image	11.9ms / 15.8ms	—	—
Resizing	10x10 input = 78.9ms 100x100 input = 1.86 s 600x600 input 56.26 s ... 1280x720 = 2min 22 s	1.1 us (from after last digit input) / 16.6 ms	99%
Box control	(infinite/blocking)	(none/very small)	100%
Box drawing	8.3 ms	(none/very small)	100%


References

[1] H. Beniwal, "Handwritten digit recognition using machine learning," *Medium*, Apr. 5, 2020. [Online]. Available: <https://medium.com/@himanshubeniwal/handwritten-digit-recognition-using-machine-learning-ad30562a9b64>.

[2] S. Gupta, "Understanding how K-means Clustering Works (a detailed guide)," *Medium*, <https://levelup.gitconnected.com/understanding-how-k-means-clustering-works-a-detailed-guide-9a2f8009a279>.

Q&A Period

Demo Setup



ML Block - K-Means

- K means estimated at 235K Ops to keep up with 15 fps.

$$\text{Confidence_N} = |(\text{IN} - \text{REF_N})|$$

(computed for each N = 0-9)

